

## Rust Lab 09 — Data Implementation & Memory (Revised)

Date: 3/9/2025

### Introduction

In this lab you will explore how Rust manages memory at a low level, including object sizes, alignment, enums, vectors, byte representation, and UTF-8 strings. Each lab exercise includes detailed tasks and spaces for observations. Answer all questions and be ready to explain your results to the TA.

### Lab A — Sizes, Alignment & #[repr]

In this exercise, you will learn how struct field ordering and alignment affect memory usage. Rust may insert padding bytes to ensure proper alignment, and different representations (`repr`) change layout guarantees.

#### Learning objectives:

- Use `std::mem::size_of`, `size_of_val`, and `align_of`.
- Observe padding and field-order effects.
- Compare `#[repr(Rust)]` vs `#[repr(C)]`.

#### Tasks:

- 1) Define the following:

```
struct S1 { a: u8, b: u64, c: u8 }
#[repr(C)] struct S2 { a: u8, b: u64, c: u8 }
struct S3 { b: u64, c: u8, a: u8 }
```

- 2) Print `size_of` and `align_of` for each.

- 3) Explain why S3 uses less memory.

- 4) Use `size_of_val(&instance)` to confirm type size.

**Expected outcome:** Students should notice padding in S1, no field reordering in S2, and tighter packing in S3 due to ordering.

#### Observation Table:

S1 size= \_\_\_\_\_ align= \_\_\_\_\_      S2 size= \_\_\_\_\_ align= \_\_\_\_\_      S3 size= \_\_\_\_\_ align= \_\_\_\_\_

Explanation: \_\_\_\_\_

TA Check: \_\_\_\_\_

### Lab B — Enums & Niche Optimization

Enums in Rust can sometimes reuse unused bit patterns, called niches, to store variant tags. This saves memory and makes `Option<T>` very efficient.

#### Learning objectives:

- Observe memory layout of `Option<T>`.
- Understand how references and `NonZero` types benefit from niche optimization.

#### Tasks:

- 1) Print `size_of::<&u8>()`, `size_of::<Option<&u8>>()`, `size_of::<u8>()`, `size_of::<Option<u8>>()`, `size_of::<NonZeroUsize>()`, `size_of::<Option<NonZeroUsize>>()`.
- 2) Compare results and explain why `Option<&u8>` is the same size as `&u8`.

**Expected outcome:** Students should see that `Option<&u8>` has the same size as `&u8` because null is reserved as `None`. For `u8`, `Option<u8>` might be larger.

#### Observation Table:

`&u8`= \_\_\_\_\_ `Option<&u8>`= \_\_\_\_\_ `u8`= \_\_\_\_\_ `Option<u8>`= \_\_\_\_\_

`NonZeroUsize`= \_\_\_\_\_ `Option<NonZeroUsize>`= \_\_\_\_\_

Explanation: \_\_\_\_\_

Student Name: \_\_\_\_\_ ID: \_\_\_\_\_ TA: \_\_\_\_\_

TA Check: \_\_\_\_\_

## Lab C — Vectors: Header vs Buffer, Growth & Preallocation

Vectors in Rust consist of a header (pointer, length, capacity) stored on the stack, and a buffer stored on the heap. When the buffer is full, Rust reallocates a larger block, usually doubling capacity.

### Learning objectives:

- Distinguish length vs capacity.
- Observe buffer growth and pointer changes.
- Use preallocation and shrink\_to\_fit.

### Tasks:

- 1) Create `Vec::new()`, push `1..=40`, and print len, cap, and pointer whenever capacity changes.
- 2) Repeat with `Vec::with_capacity(32)` and push `1..=40`.
- 3) Call `shrink_to_fit` and note capacity changes.
- 4) Convert `vec![1,2,3,4,5]` into a boxed slice and discuss use cases.

**Expected outcome:** Students should observe doubling behavior (4, 8, 16, 32...). Preallocation avoids multiple reallocations. `shrink_to_fit` may or may not reduce capacity.

### Observation Table:

Cap changes: len=\_\_\_\_, cap=\_\_\_\_, ptr=0x\_\_\_\_\_

Preallocated first cap=\_\_\_\_ growth=\_\_\_\_\_

Shrink before=\_\_\_\_ after=\_\_\_\_\_

Boxed slice: when preferable? \_\_\_\_\_

TA Check: \_\_\_\_\_

## Lab D — Bytes & Endianness: Safe vs Unsafe

Endianness describes byte order in memory. Rust provides safe functions to view integers as byte arrays. Unsafe methods like `from_raw_parts` let us view structs as bytes directly, but this can be risky.

### Learning objectives:

- Understand endianness.
- Practice safe conversions with `to_*_bytes`.
- Explore unsafe zero-copy memory views.
- Discuss risks of unsafe methods.

### Tasks:

- 1) Print `to_ne_bytes`, `to_be_bytes`, `to_le_bytes` for `0x11223344`.

- 2) Define struct `Data { i: u32, c: char, b: bool }` and view bytes:

```
#[repr(C)]  
struct Data { i: u32, c: char, b: bool }
```

(a) Safe: convert fields individually. (b) Unsafe: use `from_raw_parts` on the struct.

- 3) Compare results and explain risks of unsafe method.

**Expected outcome:** Students see that output depends on machine endianness. The `#[repr(C)]` annotation ensures consistent field ordering across compilations. Unsafe gives raw bytes including padding. They should mention risks like undefined behavior, layout instability, aliasing.

### Observation Table:

ne=\_\_\_\_ be=\_\_\_\_ le=\_\_\_\_

Raw bytes (unsafe)=\_\_\_\_\_

Safety discussion: \_\_\_\_\_

TA Check: \_\_\_\_\_

Student Name: \_\_\_\_\_ ID: \_\_\_\_\_ TA: \_\_\_\_\_

## Lab E — Strings & UTF-8

Strings in Rust are UTF-8 encoded. Indexing by bytes can cut characters, causing runtime panics. Rust enforces safety by preventing direct indexing. Use `chars()` or `graphemes` for correct iteration.

### Learning objectives:

- Work with bytes, chars, and graphemes.
- Avoid invalid slicing.
- Understand graphemes for user-visible characters.

### Tasks:

- 1) For `s="Rust 🚀 café á"`: print number of bytes, chars, and graphemes.
- 2) Show valid slice, and attempt invalid slice causing panic. **Example byte indices for `s="Rust 🚀 café á"`:**
  - Valid: `&s[0..4]` (gets "Rust")
  - Invalid: `&s[5..6]` (cuts through 🚀 emoji at byte 5, should panic)
  - Invalid: `&s[10..11]` (cuts through é at byte 10, should panic)
- 3) Extract first grapheme with `graphemes(true)`.

**Expected outcome:** Students should count differences: more bytes than chars, and graphemes may differ when accents/emoji are used. They should see why slicing incorrectly can panic - the emoji 🚀 takes 4 bytes (5-8), and é takes 2 bytes (10-11 for the combined character).

### Observation Table:

bytes=\_\_\_\_ chars=\_\_\_\_ graphemes=\_\_\_\_      Valid slice=\_\_\_\_ Invalid slice=\_\_\_\_      Grapheme example=\_\_\_\_

Explanation: \_\_\_\_\_

TA Check: \_\_\_\_\_

## Extension — Tiny Box<T> Stack

`Box<T>` lets us allocate on the heap. Recursive data structures like linked lists must use `Box` to avoid infinite size at compile time. This exercise builds a simple stack using `Box`.

**Task:** Implement `BoxedStack` with `Node` enum (`Cons`, `Nil`). Functions: `new`, `push`, `pop`, `peek`, `is_empty`, and `Debug`. Compare performance with `Vec`-based stack.

**Expected outcome:** Students practice `Box` usage and recursive enums. They should note that `Vec` is more efficient, but `Boxed` structures are useful for linked, recursive patterns.

TA Check: \_\_\_\_\_